



Politechnika Łódzka

Wydział Elektrotechniki, Elektroniki, Informatyki i Automatyki

Network Programming

Lecture 3
TCP Client



Instytut Informatyki Stosowanej
Politechniki Łódzkiej

Dr Radosław Wajman

Socket address structure

- ▶ The structure contains the protocol specific addressing information that is passed from the user process to the kernel and vice versa
- ▶ Each of the protocols supported by a socket implementation have their own socket address structure

`sockaddr_suffix`

Where *suffix* represents the protocol family

Ex: `sockaddr_in` – Internet/IPv4 socket address structure

`sockaddr_ipx` – IPX socket address structure

Socket address structure (IPv4)

The generic socket address structure

```
typedef struct sockaddr {  
    u_short sa_family;  
    CHAR sa_data[14];  
} SOCKADDR;
```

The internet/IPv4 socked address structure

```
typedef struct sockaddr_in {  
    short sin_family;  
    unsigned short sin_port;  
    IN_ADDR sin_addr;  
    CHAR sin_zero[8];  
} SOCKADDR_IN, *PSOCKADDR_IN;
```

```
typedef struct in_addr {  
    union {  
        struct {  
            UCHAR s_b1, s_b2, s_b3, s_b4;  
        } S_un_b;  
        struct {  
            USHORT s_w1, s_w2;  
        } S_un_w;  
        ULONG S_addr;  
    } S_un;  
} IN_ADDR, *PIN_ADDR, *LPIN_ADDR;
```

```
sockaddr_in service;  
service.sin_family = AF_INET;  
service.sin_addr.s_addr =  
    inet_addr("127.0.0.1");  
service.sin_port = htons(3370);
```

IP address <-> Domain names

```
sockaddr_in service;  
service.sin_family = AF_INET; ←  
service.sin_port = htons(3370);  
service.sin_addr.s_addr = ?????;
```

IP address in TXT format:

```
.s_addr =  
    inet_addr("127.0.0.1");
```

How to assign the TXT domain name address?
google.pl = 74.125.77.147?

#define AF_UNSPEC	0	// unspecified
#define AF_UNIX	1	// local to host (pipes, portals)
#define AF_INET	2	// internetwork: UDP, TCP, etc.
#define AF_IMPLINK	3	// arpanet imp addresses
#define AF_PUP	4	// pup protocols: e.g. BSP
#define AF_CHAOS	5	// mit CHAOS protocols
#define AF_NS	6	// XEROX NS protocols
#define AF_IPX	AF_NS	// IPX protocols: IPX, SPX, etc.
#define AF_ISO	7	// ISO protocols
#define AF_OSI	AF_ISO	// OSI is ISO
#define AF_ECMA	8	// european computer manufacturers
#define AF_DATAKIT	9	// datakit protocols
#define AF_CCITT	10	// CCITT protocols, X.25 etc
#define AF_SNA	11	// IBM SNA
#define AF_DECnet	12	// DECnet
#define AF_DLI	13	// Direct data link interface
#define AF_LAT	14	// LAT
#define AF_HYLINK	15	// NSC Hyperchannel
#define AF_APPLETALK	16	// AppleTalk
#define AF_NETBIOS	17	// NetBios-style addresses
#define AF_VOICEVIEW	18	// VoiceView
#define AF_FIREFOX	19	// Protocols from Firefox
#define AF_UNKNOWN1	20	// Somebody is using this!
#define AF_BAN	21	// Banyan
#define AF_ATM	22	// Native ATM Services
#define AF_INET6	23	// Internetwork Version 6
#define AF_CLUSTER	24	// Microsoft Wolfpack
#define AF_12844	25	// IEEE 1284.4 WG AF
#define AF_IRDA	26	// IrDA
#define AF_NETDES	28	// Network Designers OSI & gateway

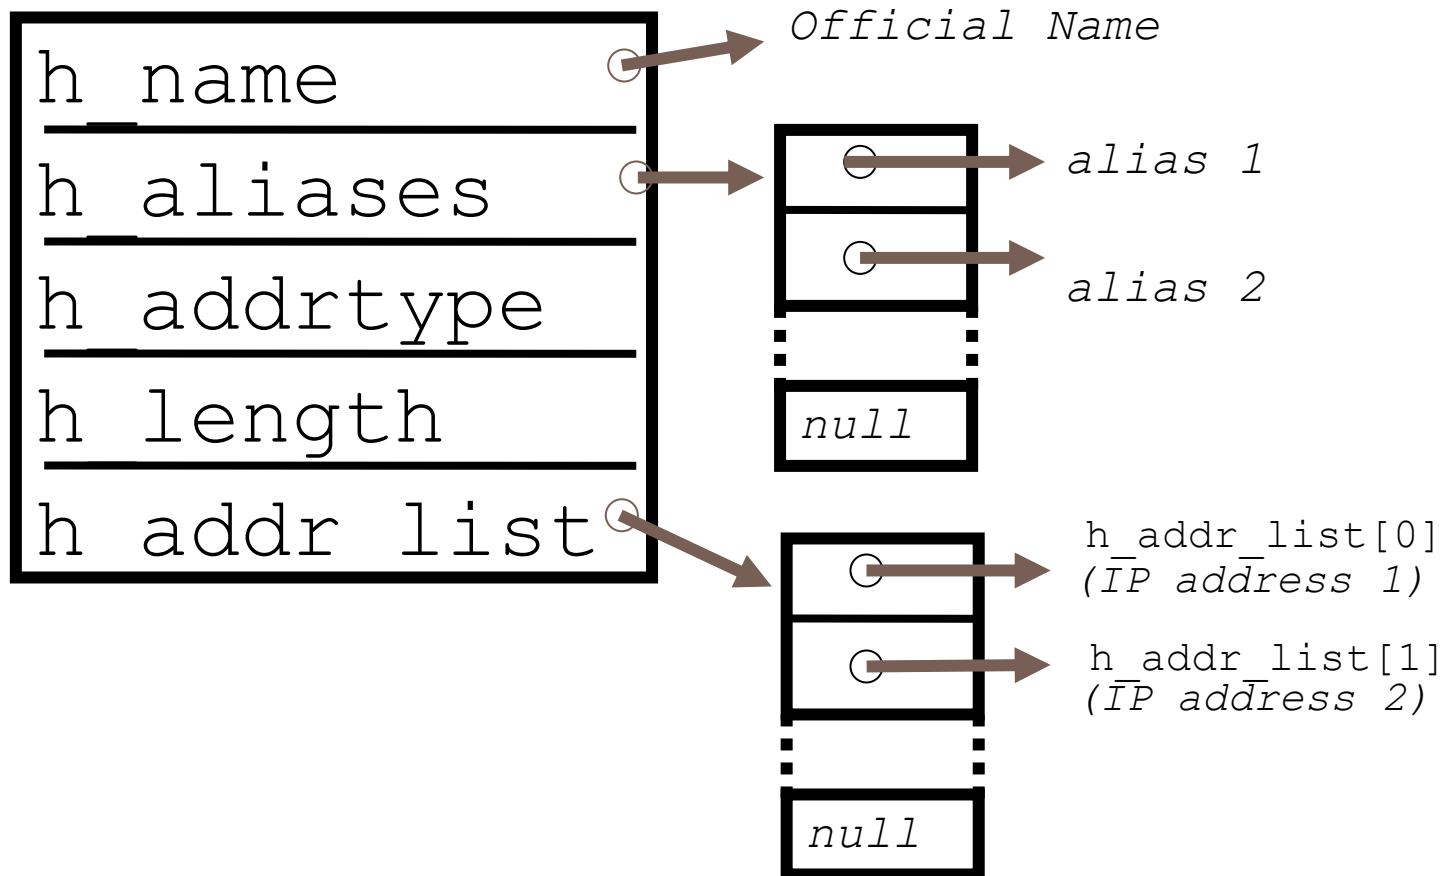
Queries for DNS

In order to send a query to DNS server and to get know the IP address of the host having its domain name the following function must be used:

```
hostent* h = gethostbyname("google.pl");  
if (h == NULL)  
{  
    printf(„error"); exit(1);  
}  
service.sin_addr = *(struct in_addr*)h->h_addr_list[0];
```

- The HOSTENT structure contains the address information,
- In the program it is not allowed to modify or deallocate it as well as any of its field.
- ONLY one copy of it is assigned to the thread,
 - Each call of **gethostbyname** or **gethostbyaddr** overwrites existing data ,

HOSTENT structure



All the IP addresses returned via the HOSTENT are in network byte order

Error handling: gethostbyname

- ▶ On error `gethostbyname` return null.
- ▶ `gethostbyname` sets the global variable `h_errno` to indicate the exact error:
 - `HOST_NOT_FOUND`
 - `TRY_AGAIN`
 - `NO_RECOVERY`
 - `NO_DATA`
 - `NO_ADDRESS`

Name/Address Conversion

- ▶ Protocol dependent DNS library functions

- `gethostbyname`
- `gethostbyaddr`

- ▶ Posix protocol *independent* functions

- `getaddrinfo()`
provides protocol-independent translation from an ANSI host name to an address
- `getnameinfo()`
provides protocol-independent name resolution from an address to an ANSI host name and from a port number to the ANSI service name
- these functions were designed to support writing code that can run on many protocols (IPv4, IPv6)

POSIX: getaddrinfo()

```
int getaddrinfo(  
    const char *hostname,  
    const char *service,  
    const struct addrinfo* hints,  
    struct addrinfo **result);
```

- ▶ `hostname` is a hostname or an address string (dotted decimal string for IP).
- ▶ `service` is a service name or a decimal port number string.
- ▶ `getaddrinfo()`
 - *provides the combined functionality of* `gethostbyname()` *and* `getservbyname()`

addrinfo structure

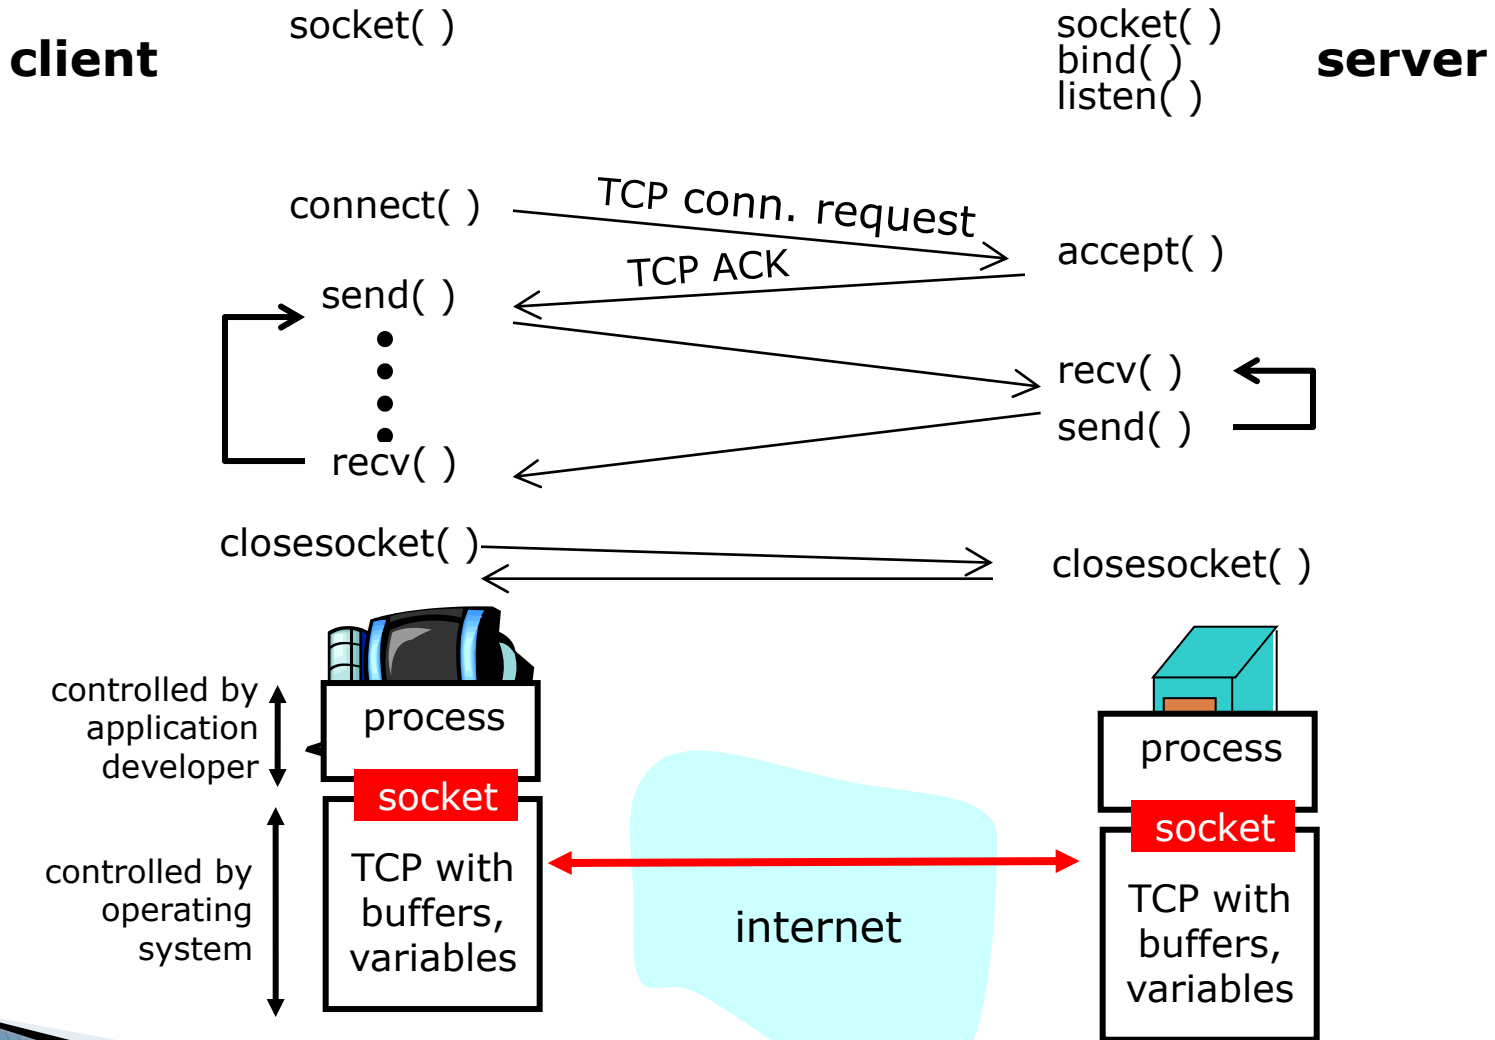
- `result` is returned with the address of a pointer to an `addrinfo` structure that is the head of a linked list.

```
struct addrinfo {  
    int      ai_flags;  
    int      ai_family;  
    int      ai_socktype;  
    int      ai_protocol;  
    size_t    ai_addrlen;  
    char      *canonicalname;  
    struct sockaddr *ai_addr;  
    struct addrinfo *ai_next;  
};
```

used by `socket()`

used by:
`bind()`
`connect()`
`sendto()`

Client+server: connection-oriented



Client: Init socket [1]

socket ► connect ► send ► recv ► closesocket

```
int socket(int family, int type, int proto);
```

Creates in the system a new socket and assigns it to the protocol.

- *family* – protocols family:
 - **AF_INET** – for IPv4,
 - **AF_INET6** – for IPv6
- *type* – socket type:
 - **SOCK_STREAM** – stream socket (TCP),
 - **SOCK_DGRAM** – datagram socket (UDP),
 - **SOCK_RAW** – raw socket,
- *proto* – protocol (for type=**SOCK_RAW**):
 - **0** – default protocol (SOCK_STREAM=TCP, SOCK_DGRAM=UDP),
- **Result: socket handle**, or:
 - **INVALID_SOCKET**, error code from *WSAGetLastError* (Windows),
 - **-1**, error code from *errno* (Unix)

Client: Init socket [2]

socket ▶ connect ▶ send ▶ recv ▶ closesocket

**In case of Windows OS,
before using socket functions,
the WinSock library must be initialized.**

```
WSADATA wsaData;  
int nCode;  
char errdesc[100];  
  
if ((nCode = WSASStartup(MAKEWORD(1, 1), &wsaData)) != 0)  
{  
    sprintf(errdesc,  
        „Error while initializing the WinSock library. Error %d“, nCode);  
    exit(1);  
}  
printf("WinSock: %s [%s]\n", wsaData.szDescription, wsaData.szSystemStatus);  
printf("MaxSockets: %d\n", wsaData.iMaxSockets);
```

Client: Init socket [3]

`socket` ▶ `connect` ▶ `send` ▶ `recv` ▶ `closesocket`

Examples of usage:

```
SOCKET sock_fd = socket(AF_INET, SOCK_STREAM, IPPROTO_TCP);
```

```
SOCKET sock_fd = socket(AF_INET, SOCK_STREAM, 0);
```

```
int sock_fd = socket(AF_INET, SOCK_STREAM, 0);
```

```
SOCKET sock_fd = socket(AF_INET, SOCK_DGRAM, 0);
```

Client: Init connection

socket ► **connect** ► send ► recv ► closesocket

```
int connect(int sock, sockaddr *rmt, int rmtlen);
```

Initializes the connection with the remote host given in *rmt*.
Processes the so called **active open**.

- ▶ *sock* – socket handle (returned by **socket**)
- ▶ *rmt* – pointer to the socket structure **sockaddr** with the remote host address, protocol and port number,
- ▶ *rmtlen* – length, in bytes, of the **sockaddr** structure for given protocol.
- ▶ **Result:** 0, or:
 - **SOCKET_ERROR**, error code from *WSAGetLastError* (Windows),
 - -1, error code from *errno* (Unix)
- ▶ **Blocking:** **connect** tries to connect with the remote host within a specified time.
- ▶ **Nonblocking:** Success or when the time is excided returns **SOCKET_ERROR** and error code = **SAEWOULDBLOCK/EWOULDBLOCK**.

```
connect(sock, (sockaddr*)&service, sizeof(sockaddr_in));
```

Client: Sending data to socket

socket ► connect ► **send** ► recv ► closesocket

or

```
int write(int sock, const char *buf, int buflen );  
int send(int sock, const char* buf, int len, int flags);
```



Sending data!

Writes data in the socket sending buffer

- ▶ *sock* – socket handle (returned by **socket** or **accept**),
- ▶ *buf* – pointer to the buffer containing data to send,
- ▶ *buflen* – count of bytes to send,
- ▶ *flags* – flags, default 0,
- ▶ **Result:** count of actually sent bytes or:
 - SOCKET_ERROR, error code from *WSAGetLastError* (Windows),
 - -1, error code from *errno* (Unix)
- ▶ **Blocking:** **send** is waiting until the sending buffer picks up *buflen* bytes
- ▶ **Nonblocking:** **send** writes the data into the buffer (no less than 1) and returns count of actually written bytes. In case of lack of buffer space **send** returns SOCKET_ERROR and the error code = WSAEWOULDBLOCK/EWOULDBLOCK.

Client: Reading data from socket

socket ► connect ► send ► **recv** ► closesocket

or
`int read(int sock, char *buf, int buflen);`
`int recv(int sock, char *buf, int buflen, int flags);`



Reads data from socket receiving buffer

- *sock* – socket handle (returned by **socket** or **accept**),
- *buf* – pointer to the buffer containing data to send,
- *buflen* – count of bytes to send,
- *flags* – flags, default 0,
- **Result:** $1 \leq \text{count of actually read bytes} \leq \text{buflen}$, or:
 - **0** – when the connection was corrupted or remotely closed,
 - **SOCKET_ERROR**, error code from *WSAGetLastError* (Windows),
 - **-1**, error code from *errno* (Unix)
- **Blocking:** **recv** waits until the minimum (default 1) bytes will arrive to the buffer
- **Nonblocking:** **recv** reads as much data as arrived (no less than 1) and returns the count of actually read bytes.
When the buffer is empty longer then the set **TIMEOUT** **recv** returns **SOCKET_ERROR** and the error code = **WSAEWOULDBLOCK/EOULDBLOCK**.

Client: Closing connection

socket ► connect ► send ► recv ► **closesocket**


or

```
int closesocket(int sock);           // windows
int close(int sock);                 // unix
```

Closes connection together with the socket.

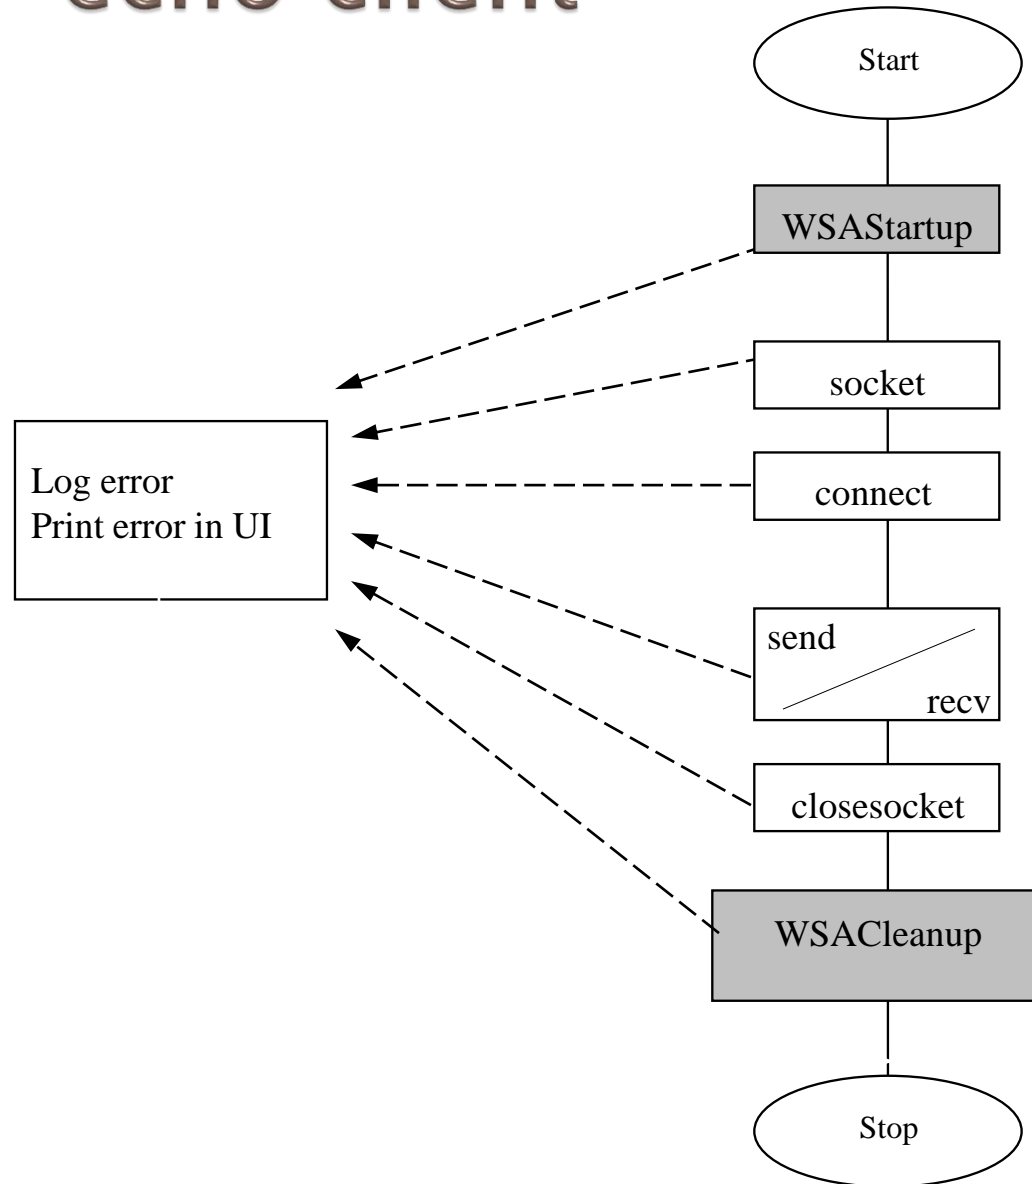
Each active operation assigned to this socket will be cancelled.

- *sock* – socket handle (returned by **socket** or **accept**),
- **Result**: **0** if socket was closed, or:
 - **SOCKET_ERROR**, error code from *WSAGetLastError* (Windows),
 - **-1**, error code from *errno* (Unix)



TCP Client

Algorithm – echo client (C++)



TCP client: C++ implementation

```
int main(int argc, char* argv[])
{
    WSADATA data;
    int result;

    result = WSAStartup(MAKEWORD(2, 0), &data);
    if(result == 0){/*handle error*/};

    SOCKET sock = socket(AF_INET, SOCK_STREAM, 0);
    if(sock != INVALID_SOCKET){/*handle error*/};

    sockaddr_in service;
    service.sin_family = AF_INET;
    service.sin_port = htons(3301);
    service.sin_addr.s_addr = inet_addr("127.0.0.1");
    result = connect(sock, (sockaddr*)&service,
                     sizeof(sockaddr_in));
    if(sock != INVALID_SOCKET){/*handle error*/};

    char str[100];
    for(int i = 0; i < 3; i++) {
        if (!read_line(sock, str))
            break;
        printf("%d: %s", i, str);
    }
    closesocket(sock);
}
```

```
bool read_line(SOCKET sock, char* line)
{
    while(true)
    {
        int result = recv(sock, line, 1, 0);
        if (result == 0 || result == SOCKET_ERROR)
            return false;
        if (*line++ == '\n')
            break;
    }
    *line = '\x0';
    return true;
}
```

The server protocol/answer

```
Data 11/10/2010\r\n
Godzina 17:53:41\r\n
Jestes klientem #1\r\n
```

TCP client: C# implementation

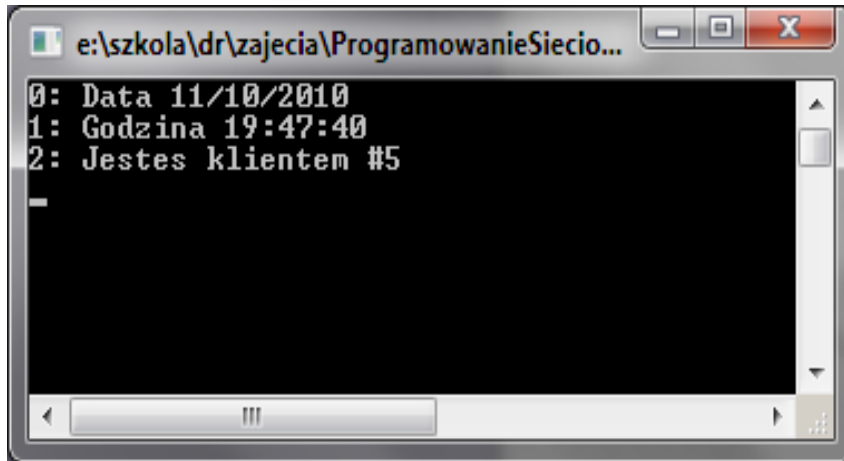
```
static void Main()
{
    Socket s = new Socket(AddressFamily.InterNetwork,
        SocketType.Stream, ProtocolType.Unspecified);
    s.Connect(new IPEndPoint(IPAddress.Parse("127.0.0.1"),
        3301));

    byte[] buffer = new byte[1024];
    int result = s.Receive(buffer);
    String time = Encoding.ASCII.GetString(buffer, 0,
        result);
    Console.WriteLine(time);
}
```

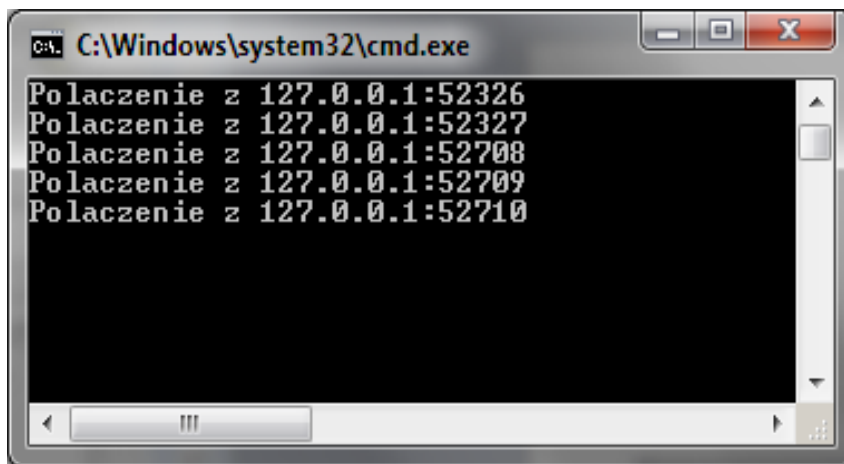
TCP client: JAVA implementation

```
1  import java.io.*;
2  import java.net.*;
3
4  class TCPClient
5  {
6      public static void main(String argv[]) throws Exception
7      {
8          String sentence;
9          Socket clientSocket = new Socket("127.0.0.1", 3301);
10         BufferedReader inFromServer = new BufferedReader(
11             new InputStreamReader(clientSocket.getInputStream()));
12         sentence = inFromServer.readLine();
13         System.out.println("FROM SERVER: " + sentence);
14         clientSocket.close();
15     }
16 }
```

TCP client and server: Tests



```
0: Data 11/10/2010
1: Godzina 19:47:40
2: Jestes klientem #5
```



```
Polaczenie z 127.0.0.1:52326
Polaczenie z 127.0.0.1:52327
Polaczenie z 127.0.0.1:52708
Polaczenie z 127.0.0.1:52709
Polaczenie z 127.0.0.1:52710
```

Server protocol / answer

```
Data 11/10/2010\r\n
Godzina 17:53:41\r\n
Jestes klientem #1\r\n
```


Error Handling

Issues and Ideas



System Errors

- ▶ In general, systems calls return a negative number to indicate an error.
 - We often want to find out what error.
 - Servers generally add this information to a log.
 - Clients generally provide some information to the user.

`extern int errno;`

- ▶ Whenever an error occurs, call *WSAGetLastError* to get error code
 - You can check error code for specific errors.
 - You can use support functions to print out or log an ASCII text error message.
- ▶ **error code** is valid only after a system call has returned an error.
 - System calls don't clear **error code** on success.
 - If you make another system call you may lose the previous value of **error code**.

Error codes

Error codes are defined in `errno.h`

<code>EAGAIN</code>	<code>EBADF</code>	<code>EACCESS</code>
<code>EBUSY</code>	<code>EINTR</code>	<code>EINVAL</code>
<code>EIO</code>	<code>ENODEV</code>	<code>EPIPE</code>

...

Support routines

- ▶ `void perror(const char *string);`
 - `stdio.h`
- ▶ `char *strerror(int errnum);`
 - `string.h`

General Strategies

- ▶ Include code to check for errors after every system call.
- ▶ Develop "wrapper functions" that do the checking for you.
- ▶ Develop layers of functions, each hides some of the error-handling details.

Example wrapper

```
int Socket( int f, int t, int p) {  
    int n;  
    if ( (n=socket(f,t,p)) < 0 ) {  
        perror("Fatal Error");  
        exit(1);  
    }  
    return(n);  
}
```

What is fatal?

- ▶ How do you know what should be a fatal error (program exits)?
 - Common sense.
 - If the program can continue – it should.but
 - if a server can't create a socket, or can't bind to it's port
 - there is no sense continuing...

Wrappers are great!

- ▶ Wrappers like those used in the text can make code much more readable.

Another approach

- ▶ Instead of simple wrapper functions, you might develop a *layered system*.
- ▶ The idea is to "hide" the `sockaddr` and error handling details behind a few custom functions or classes:

```
int tcp_client(char *server, int port);  
int tcp_server(int port);
```

Layers and Code Re-use

- ▶ Developing general functions that might be re-used in other programs is obviously "a good thing".
- ▶ Layering is beneficial even if the code is not intended to be re-used:
 - hide error-handling from "high-level" code.
 - hide other details.
 - often makes debugging easier.

The best approach to handling errors?

- ▶ There is no *best approach*.
- ▶ Do what works for you.
- ▶ Make sure you check *all* system calls for errors!
 - Not checking can lead to security problems!
 - Not checking can lead to bad grades on lab projects!